

```
In [0]: %matplotlib inline
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
from datetime import timedelta
import datetime as dt
import matplotlib.pyplot as plt
plt.rcParams['figure.figsize'] = [16, 10]
import seaborn as sns
import xgboost as xgb
from sklearn.model_selection import train_test_split
from sklearn.decomposition import PCA
from sklearn.cluster import MiniBatchKMeans
import warnings
warnings.filterwarnings('ignore')
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from scipy import stats
from scipy.stats import norm, skew #for some statistics
from sklearn.cluster import KMeans
from sklearn.cluster import KMeans
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.kernel_ridge import KernelRidge
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import KFold, cross_val_score, train_test_split
from sklearn.metrics import mean_squared_error
import random as rnd
from sklearn.metrics import mean_absolute_error
from sklearn.model_selection import train_test_split
import seaborn as sns #advanced visualization library
import requests, zipfile, io
```

```
import warnings
from datetime import datetime
warnings.filterwarnings('ignore')
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
%config InlineBackend.figure_format = 'retina' #set 'png' here when working on notebook
%matplotlib inline
```

```
In [0]: train = pd.read_csv('drive/My Drive/TAXI/train.csv')
```

```
In [0]: train.head(5)
```

Out[0]:

	id	vendor_id	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude
0	id2875421	2	2016-03-14 17:24:55	2016-03-14 17:32:30	1	-73.982155
1	id2377394	1	2016-06-12 00:43:35	2016-06-12 00:54:38	1	-73.980415
2	id3858529	2	2016-01-19 11:35:24	2016-01-19 12:10:48	1	-73.979027
3	id3504673	2	2016-04-06 19:32:31	2016-04-06 19:39:40	1	-74.010040
4	id2181028	2	2016-03-26 13:30:55	2016-03-26 13:38:10	1	-73.973053

```
In [0]: train = train[['pickup_datetime', 'dropoff_datetime', 'passenger_count',
'pickup_longitude', 'pickup_latitude', 'dropoff_longitude', 'dropoff_latitude', 'trip_duration']]
```

```
In [0]: # outliers
train=train[ train['trip_duration']<5000]
```

```
In [0]: # outliers
train=train[(train.passenger_count>=0) &(train.passenger_count<=8)]
train= train[(train.pickup_longitude>=-74.03) & (train.pickup_longitude
<=-73.77)]
train= train[(train.pickup_latitude>=40.63) & (train.pickup_latitude<=4
0.85)]
train= train[(train.dropoff_longitude>=-74.03) & (train.dropoff_longitu
de<=-73.77)]
train= train[(train.dropoff_latitude>=40.63) & (train.dropoff_latitude<
=40.85)]
train = train.dropna(how = 'any', axis = 'rows')
```

## Feature Extraction

### from Date

```
In [0]: train['pickup_datetime']=pd.to_datetime(train['pickup_datetime'],format
="%Y-%m-%d %H:%M:%S")
train.loc[:, 'pickup_weekday'] = train['pickup_datetime'].dt.weekday
train.loc[:, 'pickup_hour_weekofyear'] = train['pickup_datetime'].dt.we
ekofyear
train.loc[:, 'pickup_hour'] = train['pickup_datetime'].dt.hour
train.loc[:, 'pickup_minute'] = train['pickup_datetime'].dt.minute
train.loc[:, 'pickup_dt'] = (train['pickup_datetime'] - train['pickup_d
atetime'].min()).dt.total_seconds()
train.loc[:, 'pickup_week_hour'] = train['pickup_weekday'] * 24 + train
['pickup_hour']
```

### Distance feature

- We use PCA to transform longitude and latitude coordinates. In this case it is not about dimension reduction since we transform 2D-> 2D. The rotation could help for decision tree splits.

```
In [0]: coords = np.vstack((train[['pickup_latitude', 'pickup_longitude']].values,
                             train[['dropoff_latitude', 'dropoff_longitude']].values,
                             ))

pca = PCA().fit(coords)
train['pickup_pca0'] = pca.transform(train[['pickup_latitude', 'pickup_longitude']])[:, 0]
train['pickup_pca1'] = pca.transform(train[['pickup_latitude', 'pickup_longitude']])[:, 1]
train['dropoff_pca0'] = pca.transform(train[['dropoff_latitude', 'dropoff_longitude']])[:, 0]
train['dropoff_pca1'] = pca.transform(train[['dropoff_latitude', 'dropoff_longitude']])[:, 1]
```

- Distance

Let's calculate the distance (km) between pickup and dropoff points. Currently Haversine is used, geopy has another heuristics (vincenty() or great\_circle()) if you prefer. The cabs are not flying and we are in New York so we could check the Manhattan (L1) distance too :)

pd.DataFrame.apply() would be too slow so the haversine function is rewritten to handle arrays. We extract the middle of the path as a feature as well.

```
In [0]: def haversine_array(lat1, lng1, lat2, lng2):
    lat1, lng1, lat2, lng2 = map(np.radians, (lat1, lng1, lat2, lng2))
    AVG_EARTH_RADIUS = 6371 # in km
    lat = lat2 - lat1
    lng = lng2 - lng1
    d = np.sin(lat * 0.5) ** 2 + np.cos(lat1) * np.cos(lat2) * np.sin(lng * 0.5) ** 2
    h = 2 * AVG_EARTH_RADIUS * np.arcsin(np.sqrt(d))
    return h

def dummy_manhattan_distance(lat1, lng1, lat2, lng2):
```

```

a = haversine_array(lat1, lng1, lat1, lng2)
b = haversine_array(lat1, lng1, lat2, lng1)
return a + b

def bearing_array(lat1, lng1, lat2, lng2):
    AVG_EARTH_RADIUS = 6371 # in km
    lng_delta_rad = np.radians(lng2 - lng1)
    lat1, lng1, lat2, lng2 = map(np.radians, (lat1, lng1, lat2, lng2))
    y = np.sin(lng_delta_rad) * np.cos(lat2)
    x = np.cos(lat1) * np.sin(lat2) - np.sin(lat1) * np.cos(lat2) * np.
cos(lng_delta_rad)
    return np.degrees(np.arctan2(y, x))

train.loc[:, 'distance_haversine'] = haversine_array(train['pickup_latitude'].values, train['pickup_longitude'].values, train['dropoff_latitude'].values, train['dropoff_longitude'].values)
train.loc[:, 'distance_dummy_manhattan'] = dummy_manhattan_distance(train['pickup_latitude'].values, train['pickup_longitude'].values, train['dropoff_latitude'].values, train['dropoff_longitude'].values)
train.loc[:, 'direction'] = bearing_array(train['pickup_latitude'].values, train['pickup_longitude'].values, train['dropoff_latitude'].values, train['dropoff_longitude'].values)
train.loc[:, 'pca_manhattan'] = np.abs(train['dropoff_pca1'] - train['pickup_pca1']) + np.abs(train['dropoff_pca0'] - train['pickup_pca0'])
train.loc[:, 'center_latitude'] = (train['pickup_latitude'].values + train['dropoff_latitude'].values) / 2
train.loc[:, 'center_longitude'] = (train['pickup_longitude'].values + train['dropoff_longitude'].values) / 2

```

```


In [0]: train["diff_lat"] = abs(train.pickup_latitude - train.dropoff_latitude)
train["diff_long"] = abs(train.pickup_longitude - train.dropoff_longitude)

```

## clustering the city

Let's cluster New York City based on the pick-up and drop-off points of each taxi ride


```
In [0]: longitude = list(train.pickup_longitude) + list(train.dropoff_longitude)
latitude = list(train.pickup_latitude) + list(train.dropoff_latitude)
plt.figure(figsize = (10,10))
plt.plot(longitude,latitude,'.', alpha = 0.4, markersize = 0.05)
plt.show()
```



```
In [0]: longitude = list(train.pickup_longitude) + list(train.dropoff_longitude)
latitude = list(train.pickup_latitude) + list(train.dropoff_latitude)
loc_df = pd.DataFrame()
loc_df['longitude'] = longitude
loc_df['latitude'] = latitude
kmeans = KMeans(n_clusters=15, random_state=2, n_init = 10).fit(loc_df)
loc_df['label'] = kmeans.labels_

loc_df = loc_df.sample(200000)
plt.figure(figsize = (10,10))
for label in loc_df.label.unique():
    plt.plot(loc_df.longitude[loc_df.label == label],loc_df.latitude[loc_df.label == label],'.', alpha = 0.3, markersize = 0.3)


plt.title('Clusters of New York')
plt.show()
```



```
In [0]: fig,ax = plt.subplots(figsize = (10,10))
        for label in loc_df.label.unique():
            ax.plot(loc_df.longitude[loc_df.label == label],loc_df.latitude[loc
```



```
_df.label == label], '.', alpha = 0.4, markersize = 0.1, color = 'gray')
    ax.plot(kmeans.cluster_centers_[label,0],kmeans.cluster_centers_[la
bel,1], 'o', color = 'r')
    ax.annotate(label, (kmeans.cluster_centers_[label,0],kmeans.cluster
_centers_[label,1]), color = 'b', fontsize = 20)
ax.set_title('Cluster Centers')
plt.show()
```

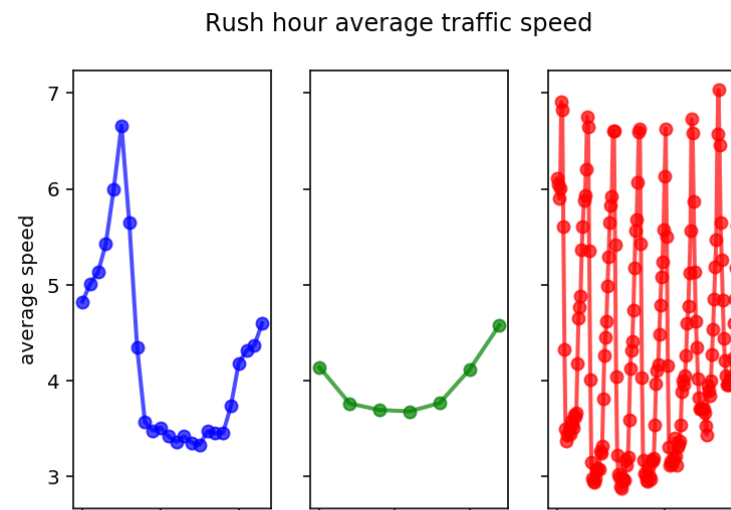


```
In [0]: train['pickup_cluster'] = kmeans.predict(train[['pickup_longitude', 'pickup_latitude']])
```

```
train['dropoff_cluster'] = kmeans.predict(train[['dropoff_longitude', 'dropoff_latitude']])
```

## Speed

```
In [0]: train.loc[:, 'avg_speed_h'] = 1000 * train['distance_haversine'] / train['trip_duration']
train.loc[:, 'avg_speed_m'] = 1000 * train['distance_dummy_manhattan'] / train['trip_duration']
fig, ax = plt.subplots(ncols=3, sharey=True)
ax[0].plot(train.groupby('pickup_hour').mean()['avg_speed_h'], 'bo-', lw=2, alpha=0.7)
ax[1].plot(train.groupby('pickup_weekday').mean()['avg_speed_h'], 'go-', lw=2, alpha=0.7)
ax[2].plot(train.groupby('pickup_week_hour').mean()['avg_speed_h'], 'ro-', lw=2, alpha=0.7)
ax[0].set_xlabel('hour')
ax[1].set_xlabel('weekday')
ax[2].set_xlabel('weekhour')
ax[0].set_ylabel('average speed')
fig.suptitle('Rush hour average traffic speed')
plt.show()
```



0      10      20      0.0      2.5      5.0      0      100  
hour                      weekday                      weekhour

## Mean speed per cluster an time

```
In [0]: new = train.groupby(['pickup_cluster', 'dropoff_cluster'])['avg_speed_h']  
        .agg({'mean_clus': 'mean'})  
        train = pd.merge(train, new, how = 'left', on=['pickup_cluster', 'dropoff_  
        cluster'])
```

```
In [0]: new = train.groupby('pickup_hour')['avg_speed_h'].agg({'mean_hour': 'mea  
n'})  
        train = pd.merge(train, new, how = 'left', on='pickup_hour')
```

```
In [0]: new = train.groupby('pickup_weekday')['avg_speed_h'].agg({'mean_weekda  
y': 'mean'})  
        train = pd.merge(train, new, how = 'left', on='pickup_weekday')
```

```
In [0]: new = train.groupby('pickup_week_hour')['avg_speed_h'].agg({'mean_weekd  
ay_hour': 'mean'})  
        train = pd.merge(train, new, how = 'left', on='pickup_week_hour')
```

## Modeling

```
In [0]: X_all = train.drop(['pickup_datetime', 'dropoff_datetime', 'avg_speed_h',  
        'avg_speed_m', 'trip_duration'], 1)  
        Y = train["trip_duration"]
```

```
In [0]: from sklearn.model_selection import KFold, cross_val_score, train_test_  
        split  
        from sklearn.preprocessing import RobustScaler, StandardScaler, MinMaxS  
        caler
```

```
X_train, val_X, Y_train, val_y = train_test_split(X_all, Y, random_state=1)
```

```
In [0]: from sklearn.metrics import mean_squared_error
        from sklearn.linear_model import ElasticNet, Lasso, BayesianRidge, LassoLarsIC
        from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
        from sklearn.kernel_ridge import KernelRidge
        from sklearn.pipeline import make_pipeline
        from sklearn.preprocessing import RobustScaler
        from sklearn.base import BaseEstimator, TransformerMixin, RegressorMixin, clone
        from sklearn.model_selection import KFold, cross_val_score, train_test_split
        from sklearn.metrics import mean_squared_error
        import xgboost as xgb
        import lightgbm as lgb
```

```
In [0]: lightgbm = lgb.LGBMRegressor(objective='regression',
                                     num_leaves=12,
                                     learning_rate=0.2,
                                     n_estimators=10000,
                                     max_bin=200,
                                     bagging_fraction=0.85,
                                     bagging_freq=5,
                                     bagging_seed=7,
                                     feature_fraction=0.2,
                                     feature_fraction_seed=7,
                                     verbose=-1,
                                     min_data_in_leaf=2,
                                     min_sum_hessian_in_leaf=11,
                                     tree_method='gpu_hist',
                                     boosting_type='gbdt',
                                     metric='rmse'
                                     )
```

```
In [0]: eval_set = [(X_train, Y_train), (val_X, val_y)]
```

```

lightgbm.fit(X_train,Y_train,eval_set=eval_set,
             verbose=50, early_stopping_rounds=100)
preds = lightgbm.predict(val_X)
val_mae2 =np.sqrt(mean_squared_error( val_y,preds))
print('valid_error',val_mae2)
preds = lightgbm.predict(X_train)
val_mae2 =np.sqrt(mean_squared_error( Y_train, preds))
print('train_error',val_mae2)

```

Training until validation scores don't improve for 100 rounds.

[50]	training's rmse: 303.421	valid_1's rmse: 303.934
[100]	training's rmse: 289.467	valid_1's rmse: 290.384
[150]	training's rmse: 283.219	valid_1's rmse: 284.471
[200]	training's rmse: 279.498	valid_1's rmse: 280.871
[250]	training's rmse: 276.886	valid_1's rmse: 278.571
[300]	training's rmse: 274.388	valid_1's rmse: 276.27
[350]	training's rmse: 272.159	valid_1's rmse: 274.2
[400]	training's rmse: 270.898	valid_1's rmse: 273.214
[450]	training's rmse: 269.479	valid_1's rmse: 272.054
[500]	training's rmse: 268.171	valid_1's rmse: 270.938
[550]	training's rmse: 267.051	valid_1's rmse: 270.028
[600]	training's rmse: 266.127	valid_1's rmse: 269.363
[650]	training's rmse: 265.2	valid_1's rmse: 268.774
[700]	training's rmse: 264.383	valid_1's rmse: 268.236
[750]	training's rmse: 263.532	valid_1's rmse: 267.624
[800]	training's rmse: 262.787	valid_1's rmse: 267.095
[850]	training's rmse: 262.104	valid_1's rmse: 266.615
[900]	training's rmse: 261.392	valid_1's rmse: 266.099
[950]	training's rmse: 260.741	valid_1's rmse: 265.703
[1000]	training's rmse: 259.985	valid_1's rmse: 265.158
[1050]	training's rmse: 259.318	valid_1's rmse: 264.685
[1100]	training's rmse: 258.737	valid_1's rmse: 264.315
[1150]	training's rmse: 258.152	valid_1's rmse: 263.911
[1200]	training's rmse: 257.561	valid_1's rmse: 263.539
[1250]	training's rmse: 257.078	valid_1's rmse: 263.291
[1300]	training's rmse: 256.588	valid_1's rmse: 263.05
[1350]	training's rmse: 256.097	valid_1's rmse: 262.743
[1400]	training's rmse: 255.657	valid_1's rmse: 262.496
[1450]	training's rmse: 255.14	valid_1's rmse: 262.198
[1500]	training's rmse: 254.676	valid_1's rmse: 262.006

```
[1550] training's rmse: 254.192      valid_1's rmse: 261.731
[1600] training's rmse: 253.758      valid_1's rmse: 261.521
[1650] training's rmse: 253.358      valid_1's rmse: 261.345
[1700] training's rmse: 252.891      valid_1's rmse: 261.069
[1750] training's rmse: 252.468      valid_1's rmse: 260.888
[1800] training's rmse: 252.065      valid_1's rmse: 260.683
[1850] training's rmse: 251.682      valid_1's rmse: 260.54
[1900] training's rmse: 251.29 valid_1's rmse: 260.367
[1950] training's rmse: 250.945      valid_1's rmse: 260.26
[2000] training's rmse: 250.604      valid_1's rmse: 260.141
[2050] training's rmse: 250.234      valid_1's rmse: 259.975
[2100] training's rmse: 249.879      valid_1's rmse: 259.846
[2150] training's rmse: 249.542      valid_1's rmse: 259.691
[2200] training's rmse: 249.201      valid_1's rmse: 259.572
[2250] training's rmse: 248.861      valid_1's rmse: 259.456
[2300] training's rmse: 248.487      valid_1's rmse: 259.282
[2350] training's rmse: 248.165      valid_1's rmse: 259.162
[2400] training's rmse: 247.83 valid_1's rmse: 259.039
[2450] training's rmse: 247.463      valid_1's rmse: 258.883
[2500] training's rmse: 247.135      valid_1's rmse: 258.745
[2550] training's rmse: 246.877      valid_1's rmse: 258.667
[2600] training's rmse: 246.577      valid_1's rmse: 258.571
[2650] training's rmse: 246.284      valid_1's rmse: 258.507
[2700] training's rmse: 246.019      valid_1's rmse: 258.403
[2750] training's rmse: 245.744      valid_1's rmse: 258.319
[2800] training's rmse: 245.454      valid_1's rmse: 258.228
[2850] training's rmse: 245.18 valid_1's rmse: 258.123
[2900] training's rmse: 244.9 valid_1's rmse: 258.064
[2950] training's rmse: 244.601      valid_1's rmse: 257.92
[3000] training's rmse: 244.348      valid_1's rmse: 257.824
[3050] training's rmse: 244.048      valid_1's rmse: 257.721
[3100] training's rmse: 243.746      valid_1's rmse: 257.603
[3150] training's rmse: 243.476      valid_1's rmse: 257.54
[3200] training's rmse: 243.194      valid_1's rmse: 257.501
[3250] training's rmse: 242.907      valid_1's rmse: 257.421
[3300] training's rmse: 242.657      valid_1's rmse: 257.391
[3350] training's rmse: 242.388      valid_1's rmse: 257.297
[3400] training's rmse: 242.156      valid_1's rmse: 257.235
[3450] training's rmse: 241.863      valid_1's rmse: 257.143
```

```
[3500] training's rmse: 241.623      valid_1's rmse: 257.082
[3550] training's rmse: 241.365      valid_1's rmse: 256.991
[3600] training's rmse: 241.1 valid_1's rmse: 256.935
[3650] training's rmse: 240.86 valid_1's rmse: 256.873
[3700] training's rmse: 240.596      valid_1's rmse: 256.817
[3750] training's rmse: 240.361      valid_1's rmse: 256.771
[3800] training's rmse: 240.117      valid_1's rmse: 256.737
[3850] training's rmse: 239.862      valid_1's rmse: 256.657
[3900] training's rmse: 239.608      valid_1's rmse: 256.577
[3950] training's rmse: 239.374      valid_1's rmse: 256.523
[4000] training's rmse: 239.164      valid_1's rmse: 256.481
[4050] training's rmse: 238.943      valid_1's rmse: 256.427
[4100] training's rmse: 238.756      valid_1's rmse: 256.395
[4150] training's rmse: 238.523      valid_1's rmse: 256.334
[4200] training's rmse: 238.302      valid_1's rmse: 256.308
[4250] training's rmse: 238.087      valid_1's rmse: 256.271
[4300] training's rmse: 237.876      valid_1's rmse: 256.226
[4350] training's rmse: 237.647      valid_1's rmse: 256.143
[4400] training's rmse: 237.457      valid_1's rmse: 256.108
[4450] training's rmse: 237.24 valid_1's rmse: 256.036
[4500] training's rmse: 237.04 valid_1's rmse: 255.998
[4550] training's rmse: 236.816      valid_1's rmse: 255.948
[4600] training's rmse: 236.598      valid_1's rmse: 255.9
[4650] training's rmse: 236.433      valid_1's rmse: 255.899
[4700] training's rmse: 236.212      valid_1's rmse: 255.844
[4750] training's rmse: 236.012      valid_1's rmse: 255.797
[4800] training's rmse: 235.806      valid_1's rmse: 255.749
[4850] training's rmse: 235.619      valid_1's rmse: 255.718
[4900] training's rmse: 235.419      valid_1's rmse: 255.656
[4950] training's rmse: 235.226      valid_1's rmse: 255.615
[5000] training's rmse: 234.973      valid_1's rmse: 255.547
[5050] training's rmse: 234.784      valid_1's rmse: 255.543
[5100] training's rmse: 234.583      valid_1's rmse: 255.51
[5150] training's rmse: 234.408      valid_1's rmse: 255.475
[5200] training's rmse: 234.24 valid_1's rmse: 255.452
[5250] training's rmse: 234.068      valid_1's rmse: 255.421
[5300] training's rmse: 233.873      valid_1's rmse: 255.385
[5350] training's rmse: 233.685      valid_1's rmse: 255.362
[5400] training's rmse: 233.519      valid_1's rmse: 255.343
```



```
[5450] training's rmse: 233.313      valid_1's rmse: 255.278
[5500] training's rmse: 233.138      valid_1's rmse: 255.278
[5550] training's rmse: 232.952      valid_1's rmse: 255.242
[5600] training's rmse: 232.769      valid_1's rmse: 255.235
[5650] training's rmse: 232.483      valid_1's rmse: 255.102
[5700] training's rmse: 232.296      valid_1's rmse: 255.077
[5750] training's rmse: 232.146      valid_1's rmse: 255.064
[5800] training's rmse: 231.982      valid_1's rmse: 255.041
[5850] training's rmse: 231.81      valid_1's rmse: 255.002
[5900] training's rmse: 231.615      valid_1's rmse: 255.001
[5950] training's rmse: 231.444      valid_1's rmse: 254.976
[6000] training's rmse: 231.246      valid_1's rmse: 254.932
[6050] training's rmse: 231.074      valid_1's rmse: 254.898
[6100] training's rmse: 230.895      valid_1's rmse: 254.857
[6150] training's rmse: 230.725      valid_1's rmse: 254.85
[6200] training's rmse: 230.55      valid_1's rmse: 254.834
[6250] training's rmse: 230.396      valid_1's rmse: 254.81
[6300] training's rmse: 230.229      valid_1's rmse: 254.815
[6350] training's rmse: 230.064      valid_1's rmse: 254.8
[6400] training's rmse: 229.917      valid_1's rmse: 254.799
[6450] training's rmse: 229.612      valid_1's rmse: 254.585
[6500] training's rmse: 229.441      valid_1's rmse: 254.572
[6550] training's rmse: 229.257      valid_1's rmse: 254.547
[6600] training's rmse: 229.083      valid_1's rmse: 254.521
[6650] training's rmse: 228.92      valid_1's rmse: 254.482
[6700] training's rmse: 228.759      valid_1's rmse: 254.451
[6750] training's rmse: 228.583      valid_1's rmse: 254.425
[6800] training's rmse: 228.394      valid_1's rmse: 254.375
[6850] training's rmse: 228.219      valid_1's rmse: 254.37
[6900] training's rmse: 228.051      valid_1's rmse: 254.349
[6950] training's rmse: 227.89      valid_1's rmse: 254.329
[7000] training's rmse: 227.734      valid_1's rmse: 254.319
[7050] training's rmse: 227.587      valid_1's rmse: 254.297
[7100] training's rmse: 227.428      valid_1's rmse: 254.298
[7150] training's rmse: 227.278      valid_1's rmse: 254.288
[7200] training's rmse: 227.118      valid_1's rmse: 254.275
[7250] training's rmse: 226.957      valid_1's rmse: 254.264
[7300] training's rmse: 226.805      valid_1's rmse: 254.261
[7350] training's rmse: 226.644      valid_1's rmse: 254.233
```

```

[7400] training's rmse: 226.479      valid_1's rmse: 254.211
[7450] training's rmse: 226.332      valid_1's rmse: 254.213
[7500] training's rmse: 226.171      valid_1's rmse: 254.186
[7550] training's rmse: 226.033      valid_1's rmse: 254.168
[7600] training's rmse: 225.877      valid_1's rmse: 254.159
[7650] training's rmse: 225.706      valid_1's rmse: 254.122
[7700] training's rmse: 225.564      valid_1's rmse: 254.108
[7750] training's rmse: 225.413      valid_1's rmse: 254.088
[7800] training's rmse: 225.269      valid_1's rmse: 254.08
[7850] training's rmse: 225.108      valid_1's rmse: 254.043
[7900] training's rmse: 224.96      valid_1's rmse: 254.048
[7950] training's rmse: 224.826      valid_1's rmse: 254.043
Early stopping, best iteration is:
[7861] training's rmse: 225.071      valid_1's rmse: 254.038
valid_error 254.0377180117667
train_error 225.07125865888986

```

```

In [0]: xgboost = xgb.XGBRegressor(colsample_bytree=0.4603, gamma=0.0468,
                                   learning_rate=0.1, max_depth=6,
                                   min_child_weight=1.7817, n_estimators=1000
00,
                                   reg_alpha=0.4640, reg_lambda=0.8571,
                                   subsample=0.5213, silent=1,
                                   random_state = 7, nthread = -1, tree_method=
'gpu_hist' )

```

```

In [0]: eval_set = [(X_train.values, Y_train), (val_X.values, val_y)]

xgboost.fit(X_train.values, Y_train, eval_set=eval_set,
            verbose=100, early_stopping_rounds=100)

preds = xgboost.predict(X_train.values)
val_mae2 = np.sqrt(mean_squared_error(Y_train, preds))
print('train_error', val_mae2)
preds = xgboost.predict(val_X.values)
val_mae2 = np.sqrt(mean_squared_error(val_y, preds))
print('vaild_error', val_mae2)

```

```

-----
XGBoostError                                Traceback (most recent call last)
<ipython-input-28-5fd006817ab8> in <module>()
      2
      3 xgboost.fit(X_train.values,Y_train,eval_set=eval_set,
----> 4             verbose=100, early_stopping_rounds=100)
      5
      6 preds = xgboost.predict(X_train.values)

/usr/local/lib/python3.6/dist-packages/xgboost/sklearn.py in fit(self,
X, y, sample_weight, eval_set, eval_metric, early_stopping_rounds, ver
bose, xgb_model, sample_weight_eval_set, callbacks)
    394             evals_result=evals_result, obj=obj
j, feval=feval,
    395             verbose_eval=verbose, xgb_model=x
gb_model,
--> 396             callbacks=callbacks)
    397
    398         if evals_result:

/usr/local/lib/python3.6/dist-packages/xgboost/training.py in train(par
ams, dtrain, num_boost_round, evals, obj, feval, maximize, early_stoppi
ng_rounds, evals_result, verbose_eval, xgb_model, callbacks, learning_r
ates)
    214             evals=evals,
    215             obj=obj, feval=feval,
--> 216             xgb_model=xgb_model, callbacks=callb
acks)
    217
    218

/usr/local/lib/python3.6/dist-packages/xgboost/training.py in _train_in
ternal(params, dtrain, num_boost_round, evals, obj, feval, xgb_model, c
allbacks)
    72         # Skip the first update if it is a recovery step.
    73         if version % 2 == 0:
--> 74             bst.update(dtrain, i, obj)

```

```

75         bst.save_rabit_checkpoint()
76         version += 1

/usr/local/lib/python3.6/dist-packages/xgboost/core.py in update(self,
dtrain, iteration, fobj)
1107         if fobj is None:
1108             _check_call(_LIB.XGBoosterUpdateOneIter(self.handle,
ctypes.c_int(iteration),
-> 1109                                     dtrain.handle))
1110         else:
1111             pred = self.predict(dtrain)

/usr/local/lib/python3.6/dist-packages/xgboost/core.py in _check_call(r
et)
174         """
175         if ret != 0:
--> 176             raise XGBoostError(py_str(_LIB.XGBGetLastError()))
177
178

```

**XGBoostError:** [12:38:59] /workspace/src/tree/updater\_gpu\_hist.cu:1407:  
Exception in gpu\_hist: NCCL failure :unhandled cuda error /workspace/s  
rc/tree/./common/device\_helpers.cuh(896)

Stack trace:

```

[bt] (0) /usr/local/lib/python3.6/dist-packages/xgboost/./lib/libxgbo
ost.so(dmlc::LogMessageFatal::~LogMessageFatal()+0x24) [0x7f45dbe6fcb4]
[bt] (1) /usr/local/lib/python3.6/dist-packages/xgboost/./lib/libxgbo
ost.so(xgboost::tree::GPUHistMakerSpecialised<xgboost::detail::Gradient
PairInternal<double> >::Update(xgboost::HostDeviceVector<xgboost::detai
l::GradientPairInternal<float> >*, xgboost::DMatrix*, std::vector<xgboo
st::RegTree*, std::allocator<xgboost::RegTree*> > const&)+0x1270) [0x7f
45dc0ab7f0]
[bt] (2) /usr/local/lib/python3.6/dist-packages/xgboost/./lib/libxgbo
ost.so(xgboost::gbm::GBTree::BoostNewTrees(xgboost::HostDeviceVector<xg
boost::detail::GradientPairInternal<float> >*, xgboost::DMatrix*, int,
std::vector<std::unique_ptr<xgboost::RegTree, std::default_delete<xgbo
ost::RegTree> >, std::allocator<std::unique_ptr<xgboost::RegTree, std::

```

```

default_delete<xgboost::RegTree> > > >*)+0xa81) [0x7f45dbef5791]
[bt] (3) /usr/local/lib/python3.6/dist-packages/xgboost/lib/libxgboost.so(xgboost::gbm::GBTree::DoBoost(xgboost::DMatrix*, xgboost::HostDeviceVector<xgboost::detail::GradientPairInternal<float> >*, xgboost::ObjFunction*)+0xd65) [0x7f45dbef6c95]
[bt] (4) /usr/local/lib/python3.6/dist-packages/xgboost/lib/libxgboost.so(xgboost::LearnerImpl::UpdateOneIter(int, xgboost::DMatrix*)+0x396) [0x7f45dbf09556]
[bt] (5) /usr/local/lib/python3.6/dist-packages/xgboost/lib/libxgboost.so(XGBoosterUpdateOneIter+0x35) [0x7f45dbe6caa5]
[bt] (6) /usr/lib/x86_64-linux-gnu/libffi.so.6(ffi_call_unix64+0x4c) [0x7f4610d09dae]
[bt] (7) /usr/lib/x86_64-linux-gnu/libffi.so.6(ffi_call+0x22f) [0x7f4610d0971f]
[bt] (8) /usr/lib/python3.6/lib-dynload/_ctypes.cpython-36m-x86_64-linux-gnu.so(_ctypes_callproc+0x2b4) [0x7f4610f1d5c4]

```

```

In [0]: from keras import optimizers
        from keras.utils import plot_model
        from keras.models import Sequential, Model
        from keras.layers.convolutional import Conv1D, MaxPooling1D
        from keras.layers import Dense, LSTM, RepeatVector, TimeDistributed, Flatten
        from sklearn.metrics import mean_squared_error
        from sklearn.model_selection import train_test_split

```

```

In [0]: epochs = 20
        batch = 256
        lr = 0.0003
        adam = optimizers.Adam(lr)
        model_mlp = Sequential()
        model_mlp.add(Dense(1024, activation='relu', input_dim=X_train.shape[1]))
        model_mlp.add(Dense(512, activation='relu'))
        model_mlp.add(Dense(256, activation='relu'))
        model_mlp.add(Dense(128, activation='relu'))
        model_mlp.add(Dense(56, activation='relu'))

```

```
model_mlp.add(Dense(1))
model_mlp.compile(loss='mse', optimizer=adam)
model_mlp.summary()
```

```
Model: "sequential_2"
```

Layer (type)	Output Shape	Param #
dense_7 (Dense)	(None, 1024)	30720
dense_8 (Dense)	(None, 512)	524800
dense_9 (Dense)	(None, 256)	131328
dense_10 (Dense)	(None, 128)	32896
dense_11 (Dense)	(None, 56)	7224
dense_12 (Dense)	(None, 1)	57

```
Total params: 727,025
Trainable params: 727,025
Non-trainable params: 0
```

```
In [0]: !pip install catboost
```

Collecting catboost

Downloading [https://files.pythonhosted.org/packages/ca/ae/aaff63662f7f5d2af7ec8d61a6f39e78ada9348e5df4f43e665ecc4bea10/catboost-0.21-cp36-none-manylinux1\\_x86\\_64.whl](https://files.pythonhosted.org/packages/ca/ae/aaff63662f7f5d2af7ec8d61a6f39e78ada9348e5df4f43e665ecc4bea10/catboost-0.21-cp36-none-manylinux1_x86_64.whl) (64.0MB)

64.0MB 59kB/s

```
Requirement already satisfied: pandas>=0.24.0 in /usr/local/lib/python
3.6/dist-packages (from catboost) (0.25.3)
```

Requirement already satisfied: matplotlib in /usr/local/lib/python3.6/dist-packages (from catboost) (3.1.3)

```
Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (from catboost) (1.12.0)
```

```
Requirement already satisfied: graphviz in /usr/local/lib/python3.6/dist-packages
```

```

t-packages (from catboost) (0.10.1)
Requirement already satisfied: numpy>=1.16.0 in /usr/local/lib/python3.6/dist-packages (from catboost) (1.17.5)
Requirement already satisfied: plotly in /usr/local/lib/python3.6/dist-packages (from catboost) (4.4.1)
Requirement already satisfied: scipy in /usr/local/lib/python3.6/dist-packages (from catboost) (1.4.1)
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.6/dist-packages (from pandas>=0.24.0->catboost) (2018.9)
Requirement already satisfied: python-dateutil>=2.6.1 in /usr/local/lib/python3.6/dist-packages (from pandas>=0.24.0->catboost) (2.6.1)
Requirement already satisfied: cyclor>=0.10 in /usr/local/lib/python3.6/dist-packages (from matplotlib->catboost) (0.10.0)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib->catboost) (2.4.6)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib->catboost) (1.1.0)
Requirement already satisfied: retrying>=1.3.3 in /usr/local/lib/python3.6/dist-packages (from plotly->catboost) (1.3.3)
Requirement already satisfied: setuptools in /usr/local/lib/python3.6/dist-packages (from kiwisolver>=1.0.1->matplotlib->catboost) (45.1.0)
Installing collected packages: catboost
Successfully installed catboost-0.21

```

```

In [0]: from catboost import CatBoostRegressor
model_cat= CatBoostRegressor(learning_rate=0.03,iterations=1500,depth=10 ,verbose = 100

)
eval_set = [(X_train, Y_train), (val_X, val_y)]

model_cat.fit(X_train.values ,Y_train)

preds = model_cat.predict(X_train.values)
val_mae2 =np.sqrt(mean_squared_error( Y_train, preds))
print('train_error',val_mae2)
preds = model_cat.predict(val_X.values)

```

```
val_mae2 = np.sqrt(mean_squared_error(val_y, preds))
print('vaild_error', val_mae2)
```

```
0:      learn: 609.8788004      total: 838ms      remaining: 20m 56s
100:    learn: 303.6896114      total: 1m 19s      remaining: 18m 19s
200:    learn: 286.5177276      total: 2m 37s      remaining: 16m 58s
300:    learn: 278.3525816      total: 3m 55s      remaining: 15m 39s
400:    learn: 272.9706923      total: 5m 14s      remaining: 14m 21s
500:    learn: 268.9150448      total: 6m 33s      remaining: 13m 4s
600:    learn: 265.5572195      total: 7m 52s      remaining: 11m 46s
700:    learn: 262.9148895      total: 9m 11s      remaining: 10m 29s
800:    learn: 260.7549750      total: 10m 30s     remaining: 9m 10s
900:    learn: 258.7712970      total: 11m 50s     remaining: 7m 52s
1000:   learn: 257.0866014      total: 13m 8s      remaining: 6m 33s
1100:   learn: 255.5272372      total: 14m 26s     remaining: 5m 13s
1200:   learn: 254.0033580      total: 15m 42s     remaining: 3m 54s
1300:   learn: 252.5841054      total: 16m 59s     remaining: 2m 35s
1400:   learn: 251.3420703      total: 18m 15s     remaining: 1m 17s
1499:   learn: 250.2308867      total: 19m 30s     remaining: 0us
train_error 250.23088669167072
vaild_error 257.4810369927262
```

```
In [0]: def ensemble_pred(X) :
        return 0.3 * xgboost.predict(X.values)+0.2*lightgbm.predict(X)+0.5*mo
        del_cat.predict(X.values)
```

## save model

```
In [0]: import torch
model_save_name = 'trip_xgb.pt'
path = F"drive/My Drive/IOT Project/{model_save_name}"
torch.save(xgboost, path)
model_save_name = 'trip_lgb.pt'
path = F"drive/My Drive/IOT Project/{model_save_name}"
torch.save(lightgbm, path)
model_save_name = 'trip_cat.pt'
```



```
path = F"drive/My Drive/IOT Project/{model_save_name}"
torch.save(model_cat, path)
```

## load the models

```
In [0]: preds = ensemble_pred(val_X)
val_X['preds'] = preds
val_X['trip_duration'] = val_y
```

```
In [0]: l = val_X.index.to_list()
```

```
In [0]: !pip install celluloid
```

```
Requirement already satisfied: celluloid in /usr/local/lib/python3.6/dist-packages (0.2.0)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.6/dist-packages (from celluloid) (3.1.3)
Requirement already satisfied: numpy>=1.11 in /usr/local/lib/python3.6/dist-packages (from matplotlib->celluloid) (1.17.5)
Requirement already satisfied: cyclur>=0.10 in /usr/local/lib/python3.6/dist-packages (from matplotlib->celluloid) (0.10.0)
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib->celluloid) (2.6.1)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib->celluloid) (2.4.6)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib->celluloid) (1.1.0)
Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (from cyclur>=0.10->matplotlib->celluloid) (1.12.0)
Requirement already satisfied: setuptools in /usr/local/lib/python3.6/dist-packages (from kiwisolver>=1.0.1->matplotlib->celluloid) (45.1.0)
```

```
In [0]: example = val_X.loc[522462]
```

```
In [0]: example
```

```
Out[0]: passenger_count      1.000000e+00
pickup_longitude      -7.399151e+01
pickup_latitude       4.074979e+01
dropoff_longitude     -7.400244e+01
dropoff_latitude      4.073362e+01
pickup_weekday        0.000000e+00
pickup_hour_weekofyear 7.000000e+00
pickup_hour           2.200000e+01
pickup_minute         2.000000e+00
pickup_dt             3.967357e+06
pickup_week_hour      2.200000e+01
pickup_pca0           -1.755646e-02
pickup_pca1           2.630611e-03
dropoff_pca0          -2.754136e-02
dropoff_pca1          1.940308e-02
distance_haversine     2.020263e+00
distance_dummy_manhattan 2.719041e+00
direction             -1.528724e+02
pca_manhattan         2.675737e-02
center_latitude       4.074170e+01
center_longitude     -7.399697e+01
diff_lat              1.617050e-02
diff_long             1.093292e-02
pickup_cluster        4.000000e+00
dropoff_cluster       1.400000e+01
mean_clus             4.000970e+00
mean_hour             4.367874e+00
mean_weekday          4.136401e+00
mean_weekday_hour     4.877751e+00
preds                 3.081132e+02
trip_duration         2.760000e+02
Name: 522462, dtype: float64
```

```
In [0]: from matplotlib import animation
from celluloid import Camera
from matplotlib.animation import FuncAnimation
x = [1 for i in range(0,20)]
fig, ax = plt.subplots(figsize=(15,8))
```

```

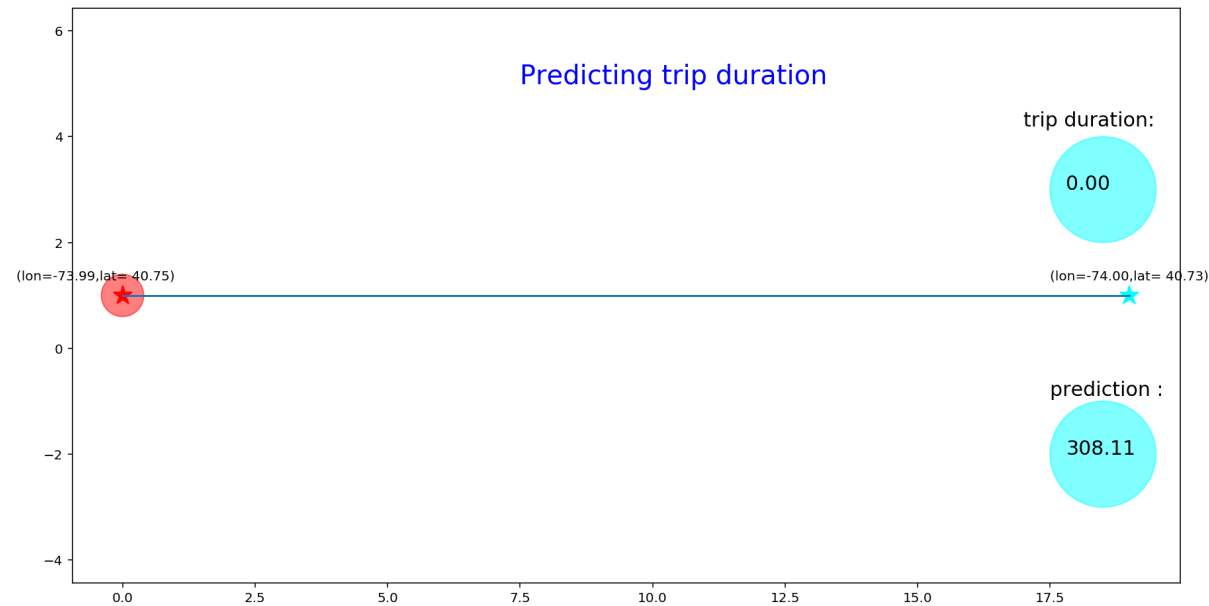
circle=plt.Circle((0,1),0.4,color='r', alpha=0.5)
ax.add_artist(circle)
ax.plot(0,1,'*',markersize=15,color='r')
ax.plot(19,1,'*',markersize=15,color='#00ffff')
ax.plot(x)
ax.text(-2,1.3,'(lon=%.2f'%example['pickup_longitude']+',lat= %.2f)' %e
xample['pickup_latitude'],{'color': 'k', 'fontsize': 10})
ax.text(17.5,1.3,'(lon=%.2f'%example['dropoff_longitude']+',lat= %.2f)'
%example['dropoff_latitude'],{'color': 'k', 'fontsize': 10})
ax.text(17,4.2,'trip duration:',{'color': 'k', 'fontsize': 15})
ax.text(17.5,-0.9,'prediction :',{'color': 'k', 'fontsize': 15})
ax.text(7.5,5,'Predicting trip duration',{'color': 'b', 'fontsize': 20
})
circle1=plt.Circle((18.5,3),1,color='#00ffff', alpha=0.5)
ax.add_artist(circle1)
circle2=plt.Circle((18.5,-2),1,color='#00ffff', alpha=0.5)
ax.add_artist(circle2)
ax.set_aspect('equal', adjustable='datalim')
text1 = ax.text(17.8,-2,'%.2f'%example['preds'],{'color': 'k', 'fontsiz
e': 15})
text2 = ax.text(17.8,3,'0',{'color': 'k', 'fontsize': 15})
n = len(np.arange(0,19,0.01))
l = np.linspace(0,example['trip_duration'],n)

def animate(i):
    circle.set_center((i,1))

    text2.set_text('%.2f'%l[int(i*100)])

animation = FuncAnimation(fig,func=animate,frames = np.arange(0,19,0.01
),interval=10 )

```



```
In [0]: animation.save('first_an.mp4', fps=100, extra_args=['-vcodec', 'libx264'])
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-1-9c1315a3749b> in <module>()
----> 1 animation.save('first_an.mp4', fps=100, extra_args=['-vcodec',
'libx264'])

NameError: name 'animation' is not defined
```

```
In [0]:
```